

ECON 425 HW3 Solutions

January 24, 2024

Problem 1

Let $f(x) = x^4 - 6x^2 + 4x + 18$. Use pen and paper to complete (i) and (ii).

- (i) Set the initial value $x^{(0)} = 1$ and the learning rate $\alpha = 0.1$. Perform three steps of gradient descent (GD) on f .
- (ii) Repeat (i) starting from the point $x^{(0)} = 0$.

Use a computer for the rest of this problem.

- (iii) Plot the function f over $x \in [-2.5, 2.5]$. How many local/global minima do you see? What are their approximate values? Can there be other local minima?
- (iv) Write your own code for S steps of GD on this function (do not use built-in or third-party GD codes).
- (v) Repeat (i) and (ii) using your code with $S = 20$ steps. Do you observe convergence in both cases?
- (vi) Now set the learning rate to $\alpha = 0.01$ and repeat (v). Explain why GD performs differently from (v).

```
[10]: # SOLUTION TO PROBLEM 1

import numpy as np
import matplotlib.pyplot as plt

S = 20 # number of steps
alpha = 0.01 # learning rate
f = lambda x: x**4 - 6*x**2 + 4*x + 18
df = lambda x: 4*x**3 - 12*x + 4

# plot the function:
x_range = np.linspace(-2.5, 2.5, 100)
plt.plot(x_range, f(x_range))
plt.grid(True)

# approximate local minima are -1.88 (also global) and 1.53

x = np.empty((S+1,))
# initial point:
```

```

x[0] = 1
print('>>> Running GD with starting value x =', x[0], 'and learning rate alpha_
↵=', alpha)

for i in range(1,S+1):
    x[i] = x[i-1] - alpha*df(x[i-1])
    print('Step', i, '  x =', x[i], '  f(x) =', f(x[i]))


```

>>> Running GD with starting value x = 1.0 and learning rate alpha = 0.01

```

Step 1  x = 1.04    f(x) = 16.84025856
Step 2  x = 1.0798054400000001    f(x) = 16.682851897094334
Step 3  x = 1.119020840068549    f(x) = 16.530861492173784
Step 4  x = 1.1572534830439267    f(x) = 16.387152244708915
Step 5  x = 1.1941304973887767    f(x) = 16.254162806484764
Step 6  x = 1.229315454262303    f(x) = 16.13373816834498
Step 7  x = 1.262522837530415    f(x) = 16.027028821639956
Step 8  x = 1.2935289442241535    f(x) = 15.93446856222769
Step 9  x = 1.3221782260494936    f(x) = 15.855828153659868
Step 10 x = 1.3483847002782723    f(x) = 15.790328788009905
Step 11 x = 1.372128707840095    f(x) = 15.736791016005489
Step 12 x = 1.3934498428239785    f(x) = 15.693793045012892
Step 13 x = 1.4124372242060033    f(x) = 15.659816283499957
Step 14 x = 1.4292183920198969    f(x) = 15.633363464226218
Step 15 x = 1.4439480114680723    f(x) = 15.613042922596875
Step 16 x = 1.4567973253950453    f(x) = 15.597619534954735
Step 17 x = 1.4679449872507762    f(x) = 15.586037367258307
Step 18 x = 1.4775696023631308    f(x) = 15.577421161028575
Step 19 x = 1.4858440513255176    f(x) = 15.571063925776986
Step 20 x = 1.4929314867366945    f(x) = 15.566406869204114

```



Problem 2

The College.csv dataset contains admissions data for a sample of 777 universities. We want to predict the number of applications received ('Apps') using the other variables in the dataset.

- (i) Let the first 600 observations be the training set and the remaining 177 observations be the test set.
- (ii) Fit the OLS regression on the training set, and report the test error obtained.

For the rest of the problem, let the penalization parameter λ vary on the 1000-point grid from 0.01 to 60.

- (iii) Fit the LASSO regression on the training set, with the penalization parameter chosen by 20-fold cross-validation. Report the test error obtained.
- (iv) Fit the ridge regression on the training set, with the penalization parameter chosen by leave-one-out cross-validation. Report the test error obtained.
- (v) Which of the three models do you prefer? Is there much difference among the test errors?

```
[11]: import matplotlib.pyplot as plt
import pandas as pd
from sklearn.linear_model import LassoCV, RidgeCV, LinearRegression
from sklearn.metrics import mean_squared_error
```

```

datapath = '/Users/franguri/Library/CloudStorage/Dropbox/_TEACHING_/Econ 425 ML_
↳UCLA Winter 2024/WEEK 3 Regularization in linear models/College.csv'
data = pd.read_csv(datapath)

# set data and grid:
Y = data['Apps']
X = data.drop('Apps', axis=1)
X['Private'] = X['Private'] == 'Yes'

Y_train = Y[:600]
X_train = X[:600]
Y_test = Y[600:]
X_test = X[600:]

alphas = np.linspace(0.01,60,1000)

# OLS:
reg = LinearRegression().fit(X_train, Y_train)
print('OLS: test MSE = ', mean_squared_error(Y_test, reg.predict(X_test)))

# LASSO:
reg = LassoCV(cv=20, random_state=0, alphas=alphas).fit(X_train, Y_train)
print('LASSO: test MSE = ', mean_squared_error(Y_test, reg.predict(X_test)))
plt.plot(reg.alphas_, np.mean(reg.mse_path_, axis=1))
plt.title('LASSO: Test MSE vs penalization parameter')
plt.show()


# RIDGE:
reg = RidgeCV(cv=None, alphas=alphas, store_cv_values=True).fit(X_train,
↳Y_train)
print('RIDGE: test MSE = ', mean_squared_error(Y_test, reg.predict(X_test)))
plt.plot(reg.alphas, np.mean(reg.cv_values_, axis=0))
plt.title('RIDGE: Test MSE vs penalization parameter')
plt.show()

```

```

OLS: test MSE = 1502077.4348215673
LASSO: test MSE = 1502356.4555399297

```



RIDGE: test MSE = 1504968.4737527145

